

Statistical learning theory

1. Learning from examples:

- Machine learning involves training from examples rather than explicit programming
- The framework focuses on supervised learning where the goal is to map the input data x into outputs y .

2. Supervised learning:

- A training set $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$
- Goal: to find an input-output relation $f(x)$ where $f(x_{new}) \cong y$

3. Data space and distribution:

- Input space $X \subseteq R^D$ output space $Y \subseteq R$.
- Data modeled by joint distribution $p(x, y)$ often factorized as $p(x, y) = p_x(x)p(y|x)$.

4. Loss function:

- A function that evaluates how much we lose if instead of real outputs we use the predicted outputs.
- $l(y_j, f(x_j)) = (y_j - f(x_j))^2$

5. Expected Loss (or expected risk):

- Is defined as the average loss over all possible $P(x, y)$ for every input-out pair, weighted by their probability of distribution.
- $\epsilon(f) = E[l(y, f(x))] = \int p(x, y)l(y, f(x))dxdy$

6. Learning Algorithms and Generalization:

- A learning algorithm computes an estimator f_s from training data S .
- Generalization: f_s should perform well on unseen data.
- Excess Expected Risk: $E(f_s) - E(f_s^*)$ Ideally minimized

7. Consistency:

- A learning algorithm is consistent, if as the number of training sample m approaches infinity, the expected risk between empirical and ideal model reaches zero.
- $\lim_{m \rightarrow \infty} E[\epsilon(f_m) - \epsilon(f^*)] = 0$

8. Fitting and stability:

- Fitting: The model should fit the training data well.
- Stability: The model shouldn't change drastically with slight variations of data.
- Overfitting: The model is too sensitive to noise (good fit, low stability)
- Oversmoothing: The model is bias (High stability, poor fit)

9. Regularization:

- a. Balance fitting and stability
- b. Sample Complexity: the minimum data size $n(\epsilon)$ needed to achieve error bound ϵ

Sample Questions:

1. What is the primary goal of supervised learning?

The primary goal of supervised learning is to find an input-output relationship $f(x)$ such that $f(x_{new}) \cong y$ for unseen data, based on the given training set $S = \{(x_1, y_1), \dots (x_n, y_n)\}$

2. Why is it necessary to model the data distribution $p(x, y)$ in statistical learning?

Modeling $p(x, y)$ accounts for uncertainty in data. The joint distribution factorizes as $p(x, y) = p_X(x)p(y|x)$, where $p_X(x)$ represents the input space variability and $p(y|x)$ accounts for noise in the outputs.

3. Explain the significance of loss function $l(y_j, f(x_j))$ in statistical learning.

The loss function quantifies the cost of predicting $f(x)$ instead of the true value y . It guides learning by providing a measure of error that the learning algorithm minimizes to estimate the target function.

4. Derive the expected risk $E(f)$ and explain its components.

$$E(f) = E[l(y, f(x))] = \int p(x, y)l(y, f(x))dxdy$$

- $p(x, y)$: joint distribution of data
- $l(y, f(x))$: Point-wise error measure
- $E(f)$: Averages the error over all possible pairs, capturing both past and future error

5. Why is the target function f^* uncomputable in practice?

The target function f^* minimizes the expected risk $E(f)$, but it depends on the true data distribution $p(x, y)$, which is unknown in practice.

6. What's the role of training set S in training?

A training set S provides samples (x_i, y_i) from $p(x, y)$. A learning algorithm uses S to compute an estimation f_S that approximates the target function f^* .

7. Define and interpret the excess expected risk $E(f_S) - E(f^*)$.

Excess expected risk quantifies the difference between the expected risk of the estimation f_S and the target function f^* . It measures how much worse f_S performs compared to the optimal model f^* .

8. What's the definition of consistency in learning algorithms?

A learning algorithm is consistent if the sample size $n \rightarrow \infty$, the expected excess risk $E_S[E(f_S) - E(f^*)] \rightarrow 0$. This implies the algorithm performance converges to the optimal model with sufficient data.

9. How does regularization help balance fitting and stability?

Regularization introduces a parameter that penalizes overfitting while ensuring the model remains stable. It prevents the model from fitting noise in the data while controlling the complexity of f_S .

10. What is sample complexity and why is it important?

A sample complexity $n(\epsilon)$ is the minimum number of samples required to achieve an error bound ϵ . It provides insight on how much data is needed for reliable learning.

11. Differentiate between overfitting and oversmoothing.

- Overfitting: the model fits the training data too well but lacks stability, leading to poor generalization.
- Oversmoothing: The model prioritizes stability but fails to capture important features, resulting in underfitting.

12. Why f_S is considered random, how does this affect generalization?

f_S depends on specific training set S , which is a random sample from $p(x, y)$. This randomness introduces variability in f_S , affecting its ability to generalize to unseen data.

13. What does the noiseless classification assumption $p(1|x) = 1 - p(-1|x)$ imply?

In noiseless classification, for any input x , $p(1|x)$ is either 1 or 0. This means there is no uncertainty in the output: the model is completely confident about class assignment.

Local methods

1. Learning from local methods

- Nearest neighbors (NN):
 - Predict y^* for a new point x^* based on the nearest neighbor in $S = \{(x_i, y_i)\}$
 - $y^* = y_i$, where $j = \operatorname{argmin}_i |x^* - x_i|$
 - Computational cost $O(nD)$ where n is the dataset size and D is the dimensionality.
- K-nearest neighbors (KNN):
 - Predict y^* as the average of K nearest neighbors.
 - $y^* = \frac{1}{K} \sum y_i$
 - Computational cost $O(nD + n \log n)$
 - It's a flexible algorithm if it's provided enough data. However, it's nervous algorithm. Meaning it's sensitive to small changes in data. That's because it relies on distance between points to make prediction.
- Parzen window:
 - Assign weights to neighbors based on proximity using a similarity function $k(x, x')$.
 - $\hat{f}(x) = \frac{\sum_{i=1}^n y_i k(x, x'_i)}{\sum_{i=1}^n k(x, x'_i)}$
 - Gaussian kernel, linear kernel, linear decay, quadric decay

2. Bias-Variance Trade-off

- Bias:
 - Difference between true function $f^*(x)$ and the expected prediction.
 - High bias -> Underfitting
- Variance:
 - Sensitivity of the model to the changes in training data.
 - High variance -> Overfitting
- Tradeoff:
 - Small K : low bias, high variance
 - Large K : high bias, low variance
 - Optimal K : Minimize total error $E_K(x) = \text{bias}^2 + \text{variance} + \text{irreducible noise}(\sigma^2)$

3. Hyperparameter selection:

- Optimal hyperparameter (K^*):
 - Minimizes the expected loss: $E_K = E_S E_{x,y} [(y - \hat{f}_{S,K}(x))^2]$
 - $K^* = \operatorname{argmin}_K \epsilon_K$
 - In practice the true distribution is unknown
 - The optimal K minimizes the sum of bias and variance, balancing model's complexity.
 - Small dataset or noisy data: use large K to reduce variance.
 - Large clean dataset: smaller K can capture finer details.
- Hold-out Cross validation:
 - Split the data into training (T) and validation (V) sets.
 - Train on T and compute the loss \hat{E}_k on (V).
 - Select K that minimizes the \hat{E}_k .
- V-fold cross validation:
 - Divide the data into V non-overlapping folds.
 - Use $V-1$ folds for training and the remaining fold for validation.

- Repeat for all V folds using each fold once as the validation.
 - Average the validation errors across all folds.
- Leave one out cross validation:
 - It's the extreme version of V -fold where $V = n$ meaning we iterate through all the data points. Use all but one data point for training, validate on the excluded point.

4. Training vs. Validation Error:

- Training Error:
 - Decreases as K increases.
- Validation Error:
 - Follows a U-shape curve, reflecting underfitting and overfitting.
 - Optimal K minimizes validation error.

5. Regression Setting:

- Noisy data model: $y = f^*(x) + \delta$, where δ is noise ($E[\delta] = 0, Var(\delta) = \sigma^2$).
- Total error $E_K(x): bias^2 + variance + \sigma^2$.

Sample Questions:

1. Define a local learning method and provide an example.

Local learning methods make prediction based on near points rather than global model.

Example: K-Nearest Neighbors (KNN) which assign the average values of K nearest neighbors of a new data point (regression) or based on the majority of classes (Classification).

2. What is the difference between KNN and NN?

- NN assigns the value of the closest training sample.
- KNN takes the majority class (Classification) or mean of K -nearest neighbors (Regression). It's less sensitive compared to NN.

3. Why is the choice of distance metric important in KNN?

The choice of metrics determines how closeness is measured between points.

- Euclidian distance: Suitable for continuous numerical data.
- Manhattan distance: works well for high-dimensional sparse data.
- Hamming distance: used for categorical or binary data.

4. Explain the computational complexity of KNN.

- Brute force of $O(nD)$ for computing the distance and $O(n \log n)$ for sorting leading to
- $O(nD + n \log n)$

5. What is Parzen window, and how does it improve KNN?

- Parzen window generalizes KNN by introducing a kernel function $K(x, x')$ that assigns weights to the neighbors based on their distances.
- Closer points have more influence over predictions.

6. Name and describe two common kernel functions in Parzen windows.

- Gaussian kernel: $K(x, x') = e^{-\frac{\|x-x'\|^2}{2\sigma}}$.
- Smoothly decreases with distance.
- Linear Kernel: $k(x, x') = (1 - \frac{\|x-x'\|}{\sigma})$
- Weights decrease linearly with distance.

7. What's bias variance tradeoff?

- The total error is decomposed into: $Error = Bias^2 + Variance + Irreducible Error$
- High bias: underfitting, over simplified model (large K in KNN)
- High variance: overfitting, model captures noise (small K in KNN)

8. How does KNN's choice of K affect the bias-variance tradeoff?

- Small K: low bias, high variance (overfitting)
- Large K: high bias, low variance (underfitting)

9. What is the irreducible error in bias-variance tradeoff?

- It is the inherent noise in the data, denoted as σ^2 , which no model can eliminate.

10. Derive the formula for KNN regression prediction:

$$\hat{f}(x^*) = \frac{1}{k} \sum_{i=0}^k y_i$$

ERM and Regularized Least Squares

1. Introduction to Regularization Networks

- Regularization Networks are a class of learning algorithms designed to prevent overfitting while maintaining generalization by incorporating regularization techniques.
- Focus on Tikhonov regularization, which **balance empirical error with model complexity**.

2. Empirical Risk Minimization (ERM):

- Empirical risk: $\hat{E}(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i))$ a proxy of uncomputable expected risk $E(f)$.
- ERM minimizes empirical risk over a hypothesis space H , but rich H can lead to overfitting.

3. Hypothesis Space and Linear Models:

- Hypothesis space: $H = \{f(x) = w^t x | x \in \mathbb{R}^d\}$.
- Linear models define predictions using a weight vector w .**

4. Tikhonov Regularization:

- Objective: $\min \hat{E}[f_w] + \lambda ||w||^2$**
- $||w||^2$ is regularization term to stabilize solution. It avoids high variance by penalizing large weights this is crucial when dealing with noisy data.**
- λ is regularization parameter. It controls the trade-off between empirical error and regularization term. High λ reduces overfitting but may cause underfitting.**

5. Regularization least square:

- Loss function: $l(y, f(x)) = (y - w^t x)^2$.**
- Optimization problem: $\min \frac{1}{n} ||Y_n - X_n w||^2 + \lambda ||w||^2$.**
- Where X_n is the input matrix and Y_n is the output vector.
- The gradient of RLS loss function:

$$-\frac{2}{n} X_n^T (Y_n - X_n w)$$

6. Computational Aspects:

- Solution involves solving a linear system: $(X_n^T X_n + \lambda I)w = X_n^T Y_n$**
- Time complexity: Training $O(nd^2)$ (when $n \gg d$) Testing: $O(d)$.**

7. Addressing Offset in Linear Models:

- The offset ensures model can represent linear functions that don't pass the origin.
- Including an offset b in the model ($f(x) = w^t x + b$):**
- Augment input space: $\bar{x} = (x, 1)$, $\bar{w} = (w, b)$.
- Regularize only w not b to avoid bias towards zero offset.
- In regularization, we penalize the magnitude of weight vector to prevent overfitting. If w includes b as well, the regularization becomes $\Rightarrow ||\bar{w}||^2 = ||w||^2 + b^2$. b is also penalized here, and the function will be biased towards $b = 0$**

8. Logistic Regression with Regularization:

- Logistic Loss function: $l(y, f_w(x)) = \log(1 + e^{-y w^t x})$.**
- Optimization Via Gradient Descent: $w_{t+1} = w_t - \gamma \nabla (\hat{E}(f_w) + \lambda ||w||^2)$**
- The Gradient of Logistic loss function:**

$$\nabla \hat{E}(f_w) = \frac{1}{n} \sum_{i=1}^n x_i \frac{-y_i}{1 + e^{y_i x_i^T w_{t-1}}}$$

- Probabilistic Interpretation: $p(1|x) = \frac{e^{w^T x}}{1 + e^{w^T x}}$
 - Logistic Regression is not just a classification problem. It also provides a probabilistic framework for predictions. So, the output of Logistic Regression can be interpreted as the probability of a data point belonging to a particular data.
9. Support Vector Machine (SVM):
- Hinge Loss: $l(y, f_w(x)) = \max(0, 1 - yw^T x)$.

$$w^* = \max_{w \in \mathbb{R}^d} \min_{1 \leq i \leq n} d(x_i, w)^2$$

$$w^* = \max_{w \in \mathbb{R}^d} \min_{1 \leq i \leq n} \frac{|x_i| - (x_i^T w)^2}{|w|^2}$$

- The optimization problem: $\min ||w||^2 + C \frac{1}{n} \sum_{i=1}^n \xi_i$, subject to $y_i w^T x_i \geq 1 - \xi_i$
 - SVM maximizes the margin between classes while allowing for some classification errors controls by C .
10. Maximum Margin Classifier
- Maximize the margin distance between the separating hyperplane and the closest data points.
 - Margin: $m_i = y_i w^T x_i$
 - Select a hyperplane equidistance from the closest points of each class.
11. Dual Formulation of SVM
- An alternative way to solve SVM optimization problem. Instead of directly optimizing the hyperplane, in the primal space, the dual formulation forces n Lagrange multipliers associated with the constraints
 - Solution is expressed as: $w = \sum_{i=1}^n \alpha_i y_i x_i$
 - Where α_i Lagrange multipliers are obtained from quadratic programming problem.
 - Sparsity: Only SVM has non-zero α
12. General Framework:
- Regularized Empirical Risk Minimized (ERM): $\min \frac{1}{n} \sum_{i=1}^n l(y_i w^T x_i) + \lambda ||w||^2$

Sample Questions:

1. What's Empirical Risk Minimization (ERM):

ERM is a framework for designing learning algorithms by minimizing empirical risk.

$$\hat{E}(f) = \frac{1}{n} \sum_{i=1}^n l(y_i, f(x_i))$$

Where $l(y, f(x))$ is the loss function, n is the number of samples. It acts as a proxy for minimizing the expected risk $E(f)$, Which is not computable due to the unknown data distribution $p(x, y)$.

2. Why is the expected risk $E(f)$ uncomputable?

The expected risk $E(f) = E[l(y, f(x))] = \int p(x, y) l(y, f(x)) dx dy$ depends on the true data distribution $p(x, y)$ which is typically unknown.

3. What is the role of hypothesis space H in ERM?

The hypothesis space H is a set of functions f over which ERM minimizes the empirical risk. It should be:

- Computationally feasible.
- Be rich enough to capture underlying patterns in the data.

4. What is an example of simple hypothesis space for linear models?

The hypothesis space for linear models is: $H = \{f(x) = w^t x | x \in \mathbb{R}^d\}$

5. What issue arises with rich hypothesis spaces?

Rich hypothesis spaces can lead to overfitting, where the model fits the training data well but performs poorly on unseen data due to capturing noise or overly complex patterns.

6. How do Tikhonov Regularization addresses overfitting?

Tikhonov Regularization adds a penalty term $\lambda ||w||^2$ to objective function, controlling model complexity ensuring stability: $\min \hat{E}[f_\omega] + \lambda ||w||^2$.

7. What is the role of regularization parameter λ ?

Regularization parameter λ balance the trade-off between fitting the data (minimizing empirical loss) and regularization (minimizing $||w||^2$).

8. What is the regularizer $||w||^2$, and why it's important?

The regularizer $||w||^2$ penalizes large weights, promoting simpler models and preventing overfitting by limiting the complexity of learned function.

9. What happens when $\lambda \rightarrow 0$ in Tikhonov Regularization?

The regularization term vanishes, and the solution minimizes empirical risk without any penalty, leading to potential overfitting.

10. How does the choice of loss function affect the solution?

Different loss functions lead to different optimization problems and learning algorithms.

- Squared loss -> Least Squares
- Logistic loss -> Logistic Regression
- Hinge loss -> SVM

11. What is the object of Regularized Least Squares (RLS)?

RLS minimizes squared error with a regularization term: $\min \frac{1}{n} \sum_{i=1}^n (y_i - w^t x_i)^2 + \lambda ||w||^2$

12. Write the RLS objective in matrix notation.

Using the input matrix X_n and output vector Y_n : $\min \frac{1}{n} ||Y_n - X_n w||^2 + \lambda ||w||^2$

13. How is the gradient of empirical risk computed for RLS?

The gradient is: $-\frac{2}{n} X_n^T (Y_n - X_n w)$

14. What is the solution to the RLS optimization problem?

The solution satisfies: $w(X_n^T X_n + \lambda n I) = X_n^T Y_n$ where λ ensures matrix is invertible.

15. Why is λ important for invertibility in RLS?

Adding λI ensures $X_n^T X_n + \lambda I$ is positive definitely, preventing issues when $X_n^T X_n$ is singular.

16. Why is Cholesky decomposition is suitable for solving RLS?

Cholesky decomposition is efficient for symmetric positive definite matrices like $X_n^T X_n + \lambda I$ making it preferred method for solving a linear system in RLS.

17. What is the computational complexity of RLS?

Training: $O(nd^2)$, assuming $n \gg d$, Testing $O(d)$.

18. How does introducing offset affect linear models?

An offset b shifts the decision boundary, allowing the model to handle data that doesn't cross the centered at origin.

The model becomes: $f(x) = w^T x + b$

19. How is offset b incorporated in RLS?

By augmenting the input vector x with an additional dimension: $\bar{x} = (x, 1)$, $\bar{w} = (w, b)$.

20. Why should the offset not be penalized in regularization?

Penalizing the offset biases the solution towards zero-offset models, which may not represent the data well.

21. How is offset b is computed in centered data?

The offset is: $b^* = \bar{y} - \bar{x}^T w^*$ where \bar{y} and \bar{x} are the means of outputs and inputs, respectively.

22. What is the advantage of centering the data in RLS?

Centering simplifies the computations by decoupling the offset b from the regularization term $\|w\|^2$

23. How is the regularization problem modified to exclude b ?

$$\min \frac{1}{n} \sum_{i=1}^n (y_i - w^T x_i - b)^2 + \lambda \|w\|^2$$

24. Why does centering data reduces the problem to standard RLS?

Centering transforms the data such that the offset b is absorbed into the mean subtraction, making the optimization problem equivalent to RLS without offset.

25. What is the Tikhonov Regularization term for centered input?

For centered input, the regularization term remains $\|w\|^2$, unaffected by the offset.

26. How does the RLS improves over standard linear regression?

The RLS adds a regularization term $\lambda \|w\|^2$, improving generalization and stability, while standard linear regression minimizes only the squared error.

27. How does the choice of λ affect RLS solutions?

- Small λ : leads to low bias but high variance increasing the risk of overfitting.
- Large λ : leads to high bias and low variance, risking the underfitting.

28. What is the objective of Regularized Logistic Regression (RLR)?

RLR minimizes the regularized empirical risk with the logistic loss:

$$\min_{w \in \mathbb{R}^d} \hat{E}(f_w) + \lambda \|w\|^2, \hat{E}(f_w) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_w(x_i)),$$

where $\ell(y, f_w(x)) = \log(1 + e^{-y f_w(x)})$.

29. Why is logistic loss function used in classification?

The logistic lost function: $\ell(y, f_w(x)) = \log(1 + e^{-y f_w(x)})$, is differentiable, convex, and provides probabilistic outputs by modeling $p(y|x)$ with a logistic function.

30. What is the probabilistic interpretation of Logistic Regression?

The logistic regression model estimates the probability of $y = 1$ given x using logistic function $p(1|x) = \frac{e^{w^T x}}{1 + e^{w^T x}}$

31. How is gradient descent applied to RLR?

Gradient descent update weights w iteratively: $w_t = w_{t-1} - \gamma \nabla (\hat{E}(f_w) + \lambda \|w\|^2)$, where the gradient of logistic loss is: $\frac{\partial \hat{E}(f_w)}{\partial w} = \frac{1}{n} \sum_{i=1}^n \frac{-y_i x_i}{1 + e^{y_i w^T x_i}}$

32. What is the role of learning rate γ in gradient descent?

The learning rate γ controls the size of step in each iteration. A small γ leads to slow convergence, while a large γ risks the overshooting the optimal solution.

33. How does the regularization improve Logistic Regression?

Regularization $\lambda \|w\|^2$ prevents overfitting by penalizing large weights, promoting generalization, especially when the dataset has noisy or high dimensional data.

34. What is the difference between RLR and unregularized Logistic Regression?

RLR adds a regularization parameter $\lambda \|w\|^2$ to the optimization objective, which controls overfitting, whereas unregularized regression only minimizes logistic loss.

35. What is the key property of Logistic function used in classification?

It maps real-valued inputs to probabilities in $(0,1)$, making it suitable for binary classification.

36. What is Support Vector Machine (SVM)?

SVMs are classification methods that finds the hyperplane that maximizes margin between two classes, ensuring robustness and good generalization.

37. What is the Hinge loss function used in SVMs?

The hinge loss is: $l(y, f_w(x)) = \max(0, 1 - yw^T x)$ it penalizes predictions that either are incorrect or withing margin.

38. Write the SVM objective with hinge loss.

The SVM optimization problem is: $\min \frac{1}{n} \sum_{i=1}^n \max(0, 1 - yw^T x) + \lambda \|w\|^2$.

39. What is margin in SVMs, and how is it defined?

The margin is the distance between hyperplane and the closest data points. For a data point (x_i, y_i) , the margin is: $m_i = y_i w^T x_i$

40. What is the geometric intuition behind the maximum margin classifier?

Classifies the training data perfectly, and maximizes the distance to the closest data points.

41. How is the SVM margin maximized mathematically?

The margin is maximized by minimizing the $\|w\|^2$ under the constraint: $y_i w^T x_i \geq 1, \forall i \in \{1, \dots, n\}$.

42. What is the primal form of the SVM optimization problem?

$$\begin{aligned} \min_{w \in \mathbb{R}^d} & \|w\|^2 + C \frac{1}{n} \sum_{i=1}^n \xi_i \\ & y_i w^T x_i \geq 1 - \xi_i, \xi_i \geq 0 \end{aligned}$$

43. What are slack variables (ξ_i) in SVM?

Slack variables allow some violation from the margin constraints, enabling SVM to handle non-linearly separable data. ξ_i measures the degree of violation for the i -th point.

44. What does parameter C controls in SVM?

The parameter C controls the trade-off between maximizing the margin and minimizing classification error:

- Large C : Focusing on minimizing error, leading to narrow margins.
- Small C : prioritizes larger margins, allowing more errors.

45. What is the dual form of SVM optimization problem?

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^n} \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j, \\ \text{subject to:} \quad & \\ & 0 \leq \alpha_i \leq C, \sum_{i=1}^n \alpha_i y_i = 0 \end{aligned}$$

46. What are support vectors in SVM?

Support vectors are data points associated with non-zero α_i in the dual solution. They lie on or within the margin and determine the hyperplane.

47. How does the dual solution determine the SVM hyperplane?

The hyperplane is defined as: $\sum_{i=1}^n \alpha_i y_i x_i$ where $\alpha_i > 0$ corresponds to support vector.

48. Why is SVM robust to outliers when C is small?

A small C allows larger margins and tolerate misclassifications, reducing the influence of outliers on the hyperplane.

49. What is the computational advantage of the dual form in SVMs?

The dual form operates in terms of α_i , requiring optimization over n variables (number of samples) instead of d variables (dimensionality of features).

50. How's the kernel trick used in SVM?

The kernel trick replaces dot product $x_i^T x_j$ with a kernel function $K(x_i, x_j)$, enabling SVMs to learn non-linear decision boundaries in higher-dimensional feature space.

51. What are common kernel functions in SVMs?

- Linear kernel: $K(x_i, x_j) = x_i^T x_j$
- Polynomial kernel: $K(x_i, x_j) = (x_i^T x_j + c)^p$
- Gaussian kernel: $K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$

52. Why does Hinge loss priorities large margins?

The Hinge loss penalizes points within the margin or misclassified points more heavily, encouraging the model to maximize the margin.

53. How does the quadratic programming formulation benefit SVMs?

Quadratic programming ensures a unique global optimum for the SVM optimization problem, as the objective function is convex.

54. What is the point-hyperplane distance in SVMs?

The distance between a point x and the hyperplane $w^T x + b = 0$ is: $d(x, w) = \frac{|w^T x + b|}{\|w\|}$

55. How is the SVM margin maximized using constraints?

The margin is maximized by solving: $\min_{w, b} \|\omega\|^2$ subject to: $y_i(w^T x_i + b) \geq 1, \forall i$.

56. What happens if data is not linearly separable in SVMs?

The constraints are relaxed by slack variables ξ_i , and a penalty term is added to the objective function: $\min_{w,b,\xi} ||\omega||^2 + C \sum_{i=1}^n \xi_i$.

57. Why are SVMs are effective for high dimensional data?

SVMs maximize the margin, which reduces overfitting even in high-dimensional spaces. The use of kernels further allows efficient computation without explicitly increasing dimensionality.

Kernel functions and feature maps

1. Extending to Non-Linear models:

- Linear models can be extended to non-linear problems by mapping input features into a higher-dimensional feature space using a feature map $\phi(x)$.
- Linear model in the feature space: $f_w(x) = w^T \phi(x)$, where $\phi: R^d \rightarrow R^D$ transforms x into higher dimensional representation.

2. Feature maps and non-linear representation:

- The feature map $\phi(x)$ allows the linear model to learn more complex, non-linear patterns in the original input space.
- Example: Mapping x to $\phi(x) = (x^2, x, 1)$ enables the model to learn quadratic relationships.

3. RLS Feature Space:

- Φ is the data matrix in feature space
- $w = (\Phi^T \Phi + \lambda n I)^{-1} \Phi^T y$
- Direct computation in high dimensions space ($D \gg d$) is infeasible due to the size of matrices like $\Phi^T \Phi$ (of size $D \times D$)
- Efficient computation is achieved by the Representer Theorem, which reduces the solution space to depend only on the data points.

4. Representer Theorem:

- Statement: for any regularized empirical risk minimization problem, the solution w^* can be expressed as: $w^* = \Phi^T c$ where $c \in R^n$ is a coefficient vector.
- This reduces the optimization problem to computation involving $n \times n$ metrics instead of $D \times D$, making it computationally feasible.

5. Kernel Trick:

- The kernel trick eliminates the need to explicitly compute $\Phi(x)$:
- Kernel function: $K(x, x') = \phi(x)^T \phi(x')$ computes the inner product of feature space without explicitly mapping x .
- Example of the function approximation: $f(x) = \sum_{i=1}^n c_i K(x, x_i)$

6. Examples of Kernel functions:

- Linear Kernel: $K(x, x') = x^T x'$
- Affine Kernel: $K(x, x') = x^T x' + \alpha^2$
- Polynomial Kernel: $K(x, x') = (x^T x' + 1)^p$
- Gaussian (RBF) Kernel: $K(x, x') = e^{-\frac{\|x-x'\|^2}{2\sigma^2}}$

7. Properties of Kernel:

A function $K: R^d \times R^d \rightarrow R$ is a valid kernel if:

- Symmetry: $K(x, x') = K(x', x)$
- Positive Definiteness: for any $\{x_1, \dots, x_n\} \subset R^d$, the matrix $K_{ij} = K(x_i, x_j)$ is positive definite.

8. Computational Efficiency using kernels:

- Using kernels, the optimization problem depends only on the kernel matrix K :

$$f(x) = \sum_{i=1}^n c_i K(x, x_i)$$

9. Applications of Kernels:

- Kernel enables the use of linear in high-dimensional feature to solve non-linear problems.
- Classification: SVM
- Regression: KRLS

10. Advantages of Kernel Trick:

- Avoids explicit computation of $\phi(x)$, reducing the computational cost.
- Handle infinite-dimensional feature space, as in the case of Gaussian Kernel.

11. Feature map is a function: $\Phi: X \rightarrow F$

- Maps input data from the input space X into a new feature space F , where learning becomes tractable.
- Feature mapping allows linear models to approximate non-linear relationships.
- The choice of $\Phi(x)$ determines the complexity and expensiveness.
- By applying transformation $\Phi(x)$ the data becomes linearly separable in higher dimension space.
- Feature maps allow nonlinear models to leverage linear methods, making computations straightforward.
- Though $\Phi(x)$ may lead to high-dimensional transformations, computation remains efficient because most methods depend only on inner products ($\Phi(x), \Phi(x')$)

Sample Questions

1. How does the feature map $\phi(x)$ extend linear models to non-linear models?

The feature map $\phi(x)$ transforms input x from the original space R^d to a higher dimensional space R^D , enabling linear models represent non-linear patterns: $f_w(x) = w^T \phi(x)$.

2. What is the computational challenge with high-dimensional feature maps?

In High-dimensional spaces ($D \gg d$) computing $\phi(x)$, $\Phi^T \Phi$ and directly solving for w becomes computationally expensive or infeasible due to the size of the matrices involved.

3. What is Representer Theorem, and why is it important?

It states that solution w^* to a regularized empirical risk minimization problem can be expressed as $w^* = \Phi^T c$ where $C \in R^n$ is a coefficient vector. **It reduces the problem to only depend on the data points, making computations feasible even in high dimensional spaces.**

4. How does the Representer Theorem simplify the computations in the feature space?

Instead of solving w in R^D , the problem is reduced to solving for c in R^n , this avoids the direct computations in high-dimensional feature space.

5. Define a kernel function.

Kernel function $K(x, x')$ computes the inner products of feature vector in high-dimensional space without explicitly computing $\phi(x)$:

$$K(x, x') = \phi(x)^T \phi(x')$$

6. How does the kernel trick work?

Kernel trick replaces the explicit the computation of $\phi(x)$ with a kernel function $K(x, x')$, allowing efficient computation of inner products in the feature space.

7. What is the mathematical expression for $f(x)$ in terms of kernel?

$$f(x) = \sum_{i=1}^n c_i K(x, x_i)$$

- Where c_i is the coefficient and $K(x, x_i)$ is the kernel function

8. What is the advantage of kernel trick in machine learning models?

The kernel trick enables learning non-linear patterns in the original input space without explicitly mapping data into high-dimensional feature space, reducing the computational cost.

9. What are the properties of a valid kernel function?

- Symmetry: $K(x, x_i) = K(x_i, x)$
- Positive definiteness: ensures the kernel can be used in optimization problems, such as solving linear systems.

10. What are linear, affine, Polynomial, and RBF kernels

- Linear Kernel: corresponding the dot product of the input vector:
$$K(x, x') = x^t x'$$
- Affine Kernel: adds a constant term to the linear Kernel:
$$K(x, x') = x^t x' + \alpha$$
- Polynomial Kernel: enables the model to represent polynomial relationships of degree p
$$K(x, x') = (x^t x' + 1)^p$$
- RBF (Gaussian) Kernel: where the σ controls the kernel's bandwidth and determines the influence of points.
$$K(x, x') = e^{-\frac{\|x - x'\|^2}{2\sigma^2}}$$

11. Why is the Gaussian kernel effective for non-linear problems?

The Gaussian kernel maps the inputs into an infinite-dimensional space, allowing the model to learn highly complex, non-linear patterns.

Neural Networks

1. Definitions:

- Deep learning: A family of machine learning methods based on artificial neural networks (ANNs) with multiple layers for hierarchical feature extraction.
- Artificial Neural Networks (ANNs): Computing systems inspired by biological neural networks, using interconnected nodes (neurons) for computations.
- Neuron Firing: Outputs $y = 1$ if weighted sum $g(X) = \sum_{i=1}^m x_i$, exceeds the threshold θ , otherwise $y = 0$.
- Activation Function:

$$f(g(X)) = \begin{cases} 1 & \text{if } g(X) > \theta \\ 0 & \text{other wise} \end{cases}$$

2. Single Layer Perceptron:

- Computes the weighted sum of input.
- Applies activation function, typically a step function

$$f(X) = \sigma(w_0 + \sum_{i=1}^m w_i x_i) = \sigma(w_0 + X^T W)$$

3. Activations Functions: Enable non-linear transformation, allowing neural networks approximate complex relationships.

- Sigmoid:
 - None-linear suitable for probabilistic interpretations.
 - Output layer of binary classification task
 - Gradient vanishing problem for very large or small data.

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- Tanh:
 - It suffers from vanishing gradient problem
 - It has better gradient properties than sigmoid
 - It is scaled version of sigmoid

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$

- ReLU:
 - Very popular and simple. Its thresholds values below 0 (**handles gradient vanishing**)
 - It allows fast convergence of the optimization function (computationally efficient)

$$f(x) = \max(0, x)$$

- Leaky ReLU:
 - It aims to fix the dying ReLU problem

$$f(x) = \begin{cases} ax & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

- Softmax:
 - Softmax allows us to represent the probability distribution over M different classes
 - Best choice for multi-class classification

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^M e^{z_k}}$$

4. Multilayer Perceptron (MLP):

- Extended single-layer perceptron by adding hidden layers.
- Each layer applies a linear transformation followed by a non-linear activation function.
- Architecture:

- The dimensionality of the hidden layer determines the width
- Number of layers determines the depth
- Each hidden unit is associated with an activation function
- $f_{\text{layer,node}}, W_{\text{layer,node(to),input(from)}}$

$$f_{1,n} = \sigma\left(\sum_{i=1}^2 w_{1,n,i} \cdot x_i\right) + b_{1,n}$$

- In an MLP, with inputs x_1, x_2 , the number of unknown parameters is:
weights(6 + 9 + 3) + bias terms(7) = 25

5. Model Capacity:

- Refers model's ability to represent complex functions. It's determined by:
 - Architecture
 - The choice of activation functions
 - The number of trainable parameters
- A model with low capacity cannot capture the patterns of data leading to underfitting.
 - **Low capacity: high bias, low variance**
- A model with high capacity, can memorize the training data and lead to overfitting.
 - **High capacity: low bias, high variance**

6. Feedforward Neural Networks

- Input signals propagate layer by layer without cycles.
- Outputs depend only on the current input and model parameter

7. Backpropagation and Training

- Training involves adjusting the weights to minimize a loss function using gradient decent.

- Computes the gradients of loss with respect to weights using the chain rule.
- Efficiently propagate errors backwards through the network.

8. Loss Function: measures the difference between the true value and predicted outputs.

- MSE for regression
- Cross-entropy loss for classification

$$W^* = \operatorname{argmin} \frac{1}{n} \sum_{i=1}^N l(F(X^i; W) y^i) = \operatorname{argmin} J(W)$$

$$w_t = w_{t-1} + \gamma \nabla J(W_{t-1})$$

9. The chain rule of derivation:

- The chain rule is the essence of DNN training, it allows to estimate gradient for high dimensional spaces from the partial derivatives with respect to each weight.

$$\frac{dF}{dx} = \frac{d}{dx} F(x) = \frac{d}{dx} f \left(g \left(h \left(u \left(v(x) \right) \right) \right) \right) = \frac{df}{dg} \cdot \frac{dg}{dh} \cdot \frac{dh}{du} \cdot \frac{du}{dv} \cdot \frac{dv}{dx}$$

10. Bagging and dropout

- Bagging: is an ensemble learning technique aimed to reduce variance and improve the generalization of the model. It works by creating multiple subsets of dataset and combining their prediction.
 - bootstrap sampling: creating multiple training dataset by sampling from the original dataset. (each subset is used to train a different dataset)
 - train independent models on each sample.
 - prediction aggregation:
 - for classification : combine prediction using the majority voting.
 - for regression : using average.
 - This is computationally very expensive; dropout is a way to approximate the same behavior.
- Dropout: At each step of the gradient descent some fraction of weights is dropped-out of each layer

11. Gradient Descent:

- Batch Gradient Descent: Computes the gradient for the entire dataset on one iteration.
 - Convergence: Ensures Convergence to the global minimum for convex cost function
 - Pros:
 - Stable and consistent convergence
 - Global perspective: gathers knowledge from all the training samples reducing the risk of noisy updates
 - Cons:
 - Computationally expensive: Slow training time on large datasets
 - Memory intensive: storing and processing entire dataset can overwhelm the memory
- $$W_{t+1} = W_t + \alpha \nabla_W J(W_t, X)$$
- $$|W_{t+1} - W_t| < \text{toll}$$
- Stochastic Gradient Descent: for each sample compute the gradient of cost function using single sample, updates the weights using the gradient, repeat the process for each sample in the dataset and several epochs.

- Pros:
 - Fast Update: each weight is updated after gradient of each sample is computed
 - Efficient memory Usage: requires to store one sample in memory in each iteration.
 - Cons:
 - Noisy Convergence: The randomness in updates introduces noise, making it harder to converge to the exact minimum.
- $$W_{t+1} = W_t + \alpha \nabla_W J(W_t, x_i)$$
- Mini-batches: Select a portion of training set of size s (mini-batch size) and update the weights after evaluating the loss function

$$W_{t+1} = W_t + \alpha \nabla_W J(W_t, x_{i:i+s})$$

Convolutional Neural Networks (CNN)

1. Introduction to CNN:

- CNN are specialized type of NN that designed for data with grid-like topology, such as:
 - 1D grids: Time-series data
 - 2D grids: Images
- They leverage two key parameters:
 - Sparse interaction: Reduced the number of connections compared to dense networks
 - Parameter sharing: Reuses the weights across different special locations.

2. Sparse Interaction:

- Unlike dense networks where each output is connected to every input, CNNs restrict connections to a small **receptive field**. With benefits of:
 - Memory Efficiency: requires fewer parameters $O(k \times n)$ than dense networks $O(m \times n)$.
 - Statistical Efficiency: Improves generalization with fewer parameters.
 - Computational Efficiency: Reduces the complexity from $O(m \times n)$ to $O(k \times n)$ where k is kernel size.

3. Parameter Sharing:

- In CNNs, the same filter (kernel) is applied across the input, detecting similar features in different regions.
- This enables efficient feature detection and reduces the number of parameters.

4. Convolution Operation:

- Convolution is a mathematical operation that involves sliding a filter(kernel) over the input to compute the feature map.
- Formula for 2D convolution: $s(i, j) = \sum_m \sum_n I(i + m, j + n)k(m, n)$. Where I is the input, k is kernel, and $s(i, j)$ is output.

5. Receptive Field:

- Each output in CNN corresponds to a specific receptive field in the input, determined by kernel size.

6. Stride and Padding:

- Stride: determines the step of filter while sliding. Larger strides results in down sampling.
- Padding: Adds zeros around the input to preserve the spatial dimensions after convolution.
- Formula for output size: $O = \frac{W-K+2P}{s} + 1$ where W is input, K is kernel, P is padding, and S is stride.

7. Pooling Layers:

- Pooling reduces special dimensions of the feature maps, improving the invariance to small shifts and reducing computational cost.
- Max pooling: Retains the maximum value in the region.
- Average pooling: Compute the average values of the region.

8. Forward and Backpropagation in CNNs:

- During forward propagation, convolution and pooling are applied to transform inputs into feature maps.
- In backpropagation:
 - Gradients are computed for each weight in the kernel by summing contributions from all positions in the output.

- This enables updating weights to minimize the loss.

9. Advantages of CNNs for Images:

- Reduced parameter count: Handle large inputs (high resolution images) efficiently.
- Feature Hierarchy: Detect low-level features (edges, texture) in initial layers and high-level features (shapes, objects) in deeper layers.
- Translation Invariance: Achieved through parameter sharing and pooling.

10. Typical CNN Architecture:

- Convolution layer: Extract features using filters
- Activation functions: Introduces non-linearity (ReLU).
- Pooling layer: Reduces dimensionality and enhances variance.
- Fully connected layer: Maps the extracted features to the output space.

11. From Feature Engineering to Feature Learning

- CNN eliminates the need to manually feature engineering by learning hierarchical features directly from data.

12. Practical Applications:

- CNNs are widely used in image classification, object detection, segmentation, and other computer vision tasks.

Bayesian Models

1. Introduction to Bayesian Learning:

- Bayesian learning focuses on estimating the probability distribution of the data and parameters, contrasting by the frequentist methods that directly estimate the quantities of interest.
- It uses prior knowledge to inform the learning process, updating the beliefs and the new data is observed.
- Bayes' theorem:

$$P(w|D) = \frac{P(D|w)P(w)}{P(D)}$$

- $P(w|D)$: Posterior (Update belief)
- $P(D|w)$: Likelihood (how well the model explains the data)
- $P(w)$: Prior (Initial belief about w)
- $P(D)$: Evidence (normalization constant)

2. Maximum Likelihood Estimation (MLE):

- MLE finds the parameters of the probabilistic model that maximize the likelihood of observed data.

$$\theta_{MLE} = \arg \max_{\theta} L(\theta|X)$$

- Example of density estimation:
 - For data $x_1 \dots x_n$ assume to follow a Gaussian distribution, the likelihood is:

$$L(\mu) \propto e^{-\sum_{i=1}^n |x_i - \mu|^2}$$

- The MLE of the mean is:

$$\mu_n = \frac{1}{n} \sum_{i=1}^n x_i$$

- Generalization: MLE can extend to multiple parameters (e.g., mean and variance) or multidimensional cases, such as multivariate Gaussian distributions.

3. Comparison of MLE and MAP:

- MLE:
 - Considers only the likelihood.
 - No prior information is incorporated.
- MAP:
 - Combines the likelihood with prior beliefs.
 - Incorporates regularization through the prior.

4. Benefits of Bayesian learning

- Uncertainty Quantification: Bayesian methods not only provide parameter estimates but also uncertainty measures, such as variance and confidence intervals.
- Flexibility: Can handle complex models and incorporate prior knowledge effectively.

5. Limitation:

- Computational cost: Bayesian methods often require integration or sampling over distributions, which can be computationally expensive for high-dimensional problems.

Sample Questions

1. What is the key difference between Maximum Likelihood Estimation (MLE) and Bayesian Learning?

- MLE estimates parameters by maximizing the likelihood function:
$$\theta_{MLE} = \arg \max_{\theta} L(\theta|X)$$
- Bayesian Learning estimates a posterior distribution over parameters using Baye's theorem:
$$P(w|D) = \frac{P(D|w)P(w)}{P(D)}$$
- MLE provides point estimate, Bayesian method incorporates prior knowledge and quantify uncertainty.

2. What are the advantages and limitations of Bayesian Learning?

- Pros:
 - Uncertainty Quantification: Bayesian methods provide a distribution over parameters instead of a single estimate.
 - Regularization: Avoids overfitting by incorporating prior knowledge.
 - Flexibility: Can handle complex models and integrate domain-specific knowledge.
- Cons:
 - Computational Cost: Bayesian inference often requires integration or approximation techniques
 - Choice of Prior: Incorrect priors can lead to biased results.
 - High-Dimensionality Issues: Exact Bayesian inference becomes infeasible for large-scale problems

3. Why is MLE sensitive to overfitting?

- MLE only maximizes the likelihood, which can fit noise in small datasets.
- No regularization is applied, leading to high variance in small sample sizes.

4. What is the role of the likelihood function in Bayesian inference?

- The likelihood function measures how well a given parameter explains the observed data.
- In Bayesian inference, the likelihood is combined with the prior to compute the posterior.

Decision Tree

1. Partition-Based Estimators

- Partition-based methods divide the input space X into disjoint sets (cells). Each cell has its own local estimate.
- Dyadic Partition
 - The input space X is divided into cubes of side length 2^{-k}
 - This ensures a hierarchical structure of partitions.

2. Piecewise Constant Estimators

- In each cell, the function is approximated by a constant.
- The goal is to minimize:

$$\min_{c_1, \dots, c_J \in \mathbb{R}} \sum_{j=1}^J \sum_{x_i \in A_j} (y_i - c_j)^2$$

- The optimal solution for each cell is:

$$\hat{c}_j = \frac{1}{n_j} \sum_{x_i \in A_j} y_i$$

- Where n_j is the number of points in each cell
- This is related to ERM, where the hypothesis space consists of piecewise constant functions.

3. Piecewise Linear Estimators

- Instead of a constant, a linear function is fit within each cell:

$$f(x) = \sum_{j=1}^J (w_j^T x + b_j) 1_{A_j}(x)$$

- The estimation problem reduces to Ordinary Least Squares (OLS) in each partition.

4. General Partition-Based Estimators

- The function in each partition can belong to any hypothesis class G (e.g., neural networks, reproducing kernel Hilbert spaces (RKHS), etc.).
- ERM is performed separately in each partition.

5. Partition Trees and Decision Trees

- Partition trees generalize the previous ideas by creating hierarchical partitions.
- Partition Tree Estimators
 - Each level q of the partition has J_q cells.
 - The function space is:

$$H_{G,q} = \{f: X \rightarrow \mathbb{R} \mid f(x) = \sum_{j=1}^{J_q} g(x) 1_{A_j}(x), g \in G\}$$

- The estimation problem remains an ERM problem within each partition.

6. Uniform vs. Adaptive Partitioning

- Uniform partitioning: The scale parameter q acts as a regularization parameter.
- Adaptive partitioning: The cells are refined dynamically based on the loss function:

$$L_A = \min_{f \in H} \frac{1}{n_A} \sum_{x_i \in A} (y_i - f_A(x_i))^2$$

- A cell is kept if $L_A \leq \tau$, otherwise, it is split.
- The threshold τ acts as a regularization parameter.

7. Decision Tree

- Decision trees use a greedy approach to partition the space.
- The partitioning criterion minimizes:

$$\min_{j,s} [L_{A_1(s,j)} + L_{A_2(s,j)}]$$
- Recursive splitting is applied to further refine the partitions.
- Alternative error measures like cross-entropy (for classification) or Gini index are used.

8. Kernel Regression

- Kernel regression shifts the focus from discrete partitions to a continuous weighting approach.

9. Nearest-Neighbor-Like Estimation

- Instead of fixed cells, each point is assigned a local weight:

$$f(\hat{x}) = \arg \min_{c \in R} \sum_{i=1}^n (y_i - c)^2 1_{(|x - x_i| \leq \tau)}$$
- This is equivalent to nearest neighbor regression, where only the closest points contribute to the estimate.

10. Nadaraya-Watson Kernel Regression

- Instead of a hard threshold, a kernel function assigns weights to nearby points:

$$f(\hat{x}) = \frac{\sum_{i=1}^n y_i k(x_i, x)}{\sum_{i=1}^n k(x_i, x)}$$
- $k(x_i, x)$ is a kernel function that determines how much weight is given to x_i when estimating $f(x)$.
- This is a weighted averaging method where closer points have higher influence.

11. Locally Linear Kernel Regression

- Instead of estimating a constant, a linear model is fitted locally:

$$f(x) = \hat{w}^T x + \hat{b}$$
- The weights are computed by solving:

$$(\hat{w}, \hat{b}) = \arg \min_{(w,b) \in R^{d+1}} \sum_{i=1}^n (y_i - w^T(x_i - x) - b)^2 k(x_i, x)$$
- This improves accuracy by allowing for local variations in slope.

12. Key Properties of Kernel Regression

- No optimization required: The solution is computed directly via weighted sums.
- Nonparametric: The function estimate adapts to the data, avoiding strong assumptions.
- Locally Adaptive: Can capture local structures in the data but does not impose global smoothness.

Sample Questions

- How do decision trees lead to overfitting, and how can it be prevented?
 - Deep trees capture noise instead of general patterns, reducing generalization ability.
 - Prevention: Pruning, Max Depth, Random Forest

2. How does Locally Weighted Regression differ from Ordinary Least Squares (OLS)?

- OLS: Finds one global linear model for the data.
- Locally Weighted Regression: Fits a linear model locally using a kernel-weighted loss function:

$$(\hat{w}, \hat{b}) = \arg \min_{(w,b) \in \mathbb{R}^{d+1}} \sum_{i=1}^n (y_i - w^T(x_i - x) - b)^2 k(x_i, x)$$

- It can model non-linearity better than OLS.

3. How does adaptive partitioning work in decision trees?

- The partitioning splits regions dynamically based on data distribution.
- A split is performed if the local loss function exceeds a threshold:

$$L_A = \min_{f \in H} \frac{1}{n_A} \sum_{x_i \in A} (y_i - f_A(x_i))^2$$

- Effect: More splits in complex areas, fewer in smooth regions.

Variable Selection

1. Introduction:

- In machine learning, prediction alone is not enough; interpretability is also crucial.
- Interpretability depends on detecting key variables that contribute to predictions.
- The goal is to select relevant variables to enhance prediction accuracy while reducing complexity.

2. Linear Models and Sparsity

- The model is defined as:

$$f_w(x) = w^T x = \sum_{i=1}^v w_i x_i$$

- x_i represents the input features (pixels, word count, ...)
- Variable selection helps identify features that significantly impact prediction
- Sparsity assumption: The best model is often sparse, meaning only a few weights w_i are nonzero.

4. High-Dimensional Statistics

- Traditional models assume $n \gg D$ (more samples than features) which result in overdetermined system.
- In modern settings, $n \ll D$ (more features than samples) which result in underdetermined system.
- Buzzwords: compressed sensing, high-dimensional statistics

5. Brute Force Approach (Hard)

- Selecting features using combinatorial search (checking all subsets of features) is computationally infeasible

$$\min_{w \in \mathbb{R}^D} \frac{1}{n} \sum_{i=1}^n (y_i - f_w(x_i))^2 + \lambda |w|_0$$

- The l_0 norm counts non-zero coefficients.
- $l_0 \rightarrow ||w||_0 = \{j | w^j \neq 0\}$
- Computational complexity grows exponentially.
- Solution: Approximation Approaches
 - Greedy methods (Matching Pursuit, Orthogonal Matching Pursuit)
 - Convex Relaxation (LASSO, Elastic Net)

6. Greedy Methods: Matching Pursuit (MP)

- Iteratively select features that are most correlated with the residual error.
- Steps:
 - Initialize: Residual $r_0 = Y$, weight vector $w_0 = [\dots 0 \dots]$, and selected feature $I_0 = \emptyset$
 - Select the variable that most correlated with the residual.

$$k = \arg \max_{j=1, \dots, D} a_j, \quad a_j = \frac{(r_{i-1}^T X^j)^2}{|X^j|^2},$$

- Update the coefficient vector w .

$$w_t = \left(\widehat{X_t^T X_t} \right)^{-1} \widehat{X_t^T Y}$$

- Update the residual.

$$r_t = \hat{Y} - \hat{X} w_t$$

7. Orthogonal Matching Pursuit (OMP)

- Instead of simple updates, OMP solves a least squares problem:

$$w_i = \arg \min_{w \in \mathbb{R}^D} \|Y_n - X_n M_i w\|^2$$
- Provides a more accurate approximation
- Works only if true solution is sparse and the features are not too correlated

8. Convex Relaxation: LASSO

- Instead of l_0 - norm, uses l_1 - norm:

$$|w|_1 = \sum_{j=1}^D |w^j|$$

- It's the sum of absolute values of weights
- The optimization problem becomes:

$$\min_{w \in \mathbb{R}^D} \frac{1}{n} \sum_{i=1}^n (y_i - f_w(x_i))^2 + \lambda |w|_1$$

- Convex can be optimized efficiently.
- Encourages sparse solutions (some weights are exactly zero).

9. Iterative Soft Thresholding Algorithm (ISTA)

- A gradient-based approach to solve LASSO:

$$w_i = S_{\lambda \gamma} \left(w_{i-1} - \frac{2\gamma}{n} X_n^T (Y_n - X_n w_{i-1}) \right)$$

- Soft thresholding function ensures sparsity:

$$S_\alpha(u) = \text{sign}(u) \cdot \max(|u| - \alpha, 0)$$
- Converges under proper step-size selection:

$$\gamma = \frac{n}{2|X_n^T X_n|}$$

10. Elastic Net: Combining LASSO and Ridge

- Problem with LASSO: When features are highly correlated, it randomly selects one.
- Solution: Elastic Net combines LASSO and Ridge regression:

$$\min_{w \in \mathbb{R}^D} \frac{1}{n} \sum_{i=1}^n (y_i - f_w(x_i))^2 + \lambda(\alpha|w|_1 + (1-\alpha)|w|_2^2)$$

- l_1 – norm Sparsity
- l_2 – norm Stability

11. ISTA for Elastic Net

- Modified iterative update:

$$w_0 = 0, w_i = S_{\lambda\alpha\gamma} \left((1 - \lambda\gamma(1-\alpha))w_{i-1} - \frac{2\gamma}{n} X_n^T (Y_n - X_n w_{i-1}) \right), i = 1, \dots, T_{max}$$

- At each iteration a non-linear soft thresholding operator is applied to the gradient step.
- The iteration should be run until a convergence criterion is met or a maximum number of iteration T_{max} is reached.
- To ensure convergence we should choose a step-size:

$$\gamma = \frac{n}{2(|X_n^T X_n| + \lambda(1-\alpha))}$$

12. Conclusion:

- Feature selection improves interpretability and reduces complexity
- Brute force search is infeasible, approximation methods are needed.
 - Greedy methods: Matching Pursuit (MP), Orthogonal Matching Pursuit (OMP)
 - Convex Relation (LASSO, Elastic Net)
- LASSO uses l_1 -norm for sparsity.
- Elastic Net blends with LASSO and Ridge for better stability.
- ISTA provides an efficient way to solve LASSO and Elastic Net.

Sample Questions

1. What's the key assumption behind sparsity-based models in machine learning?

The key assumption is that the best predictive model only relies on small subset of features. This means that most features weights in the model should be zero, leading to a sparse solution.

2. Explain why the l_0 -norm regularization is computationally infeasible.

l_0 -norm regularization requires a combinatorial search over all possible features subsets to determine which subset minimizes the loss. This result is an NP-hard problem, making it infeasible for large datasets.

3. Why l_1 -norm regularization (LASSO) preferred over l_0 -norm in practice?

l_1 -norm regularization provides a convex relaxation of the l_0 -norm making the optimization problem computationally tractable. It encourages sparsity while allowing efficient solutions using gradient-based methods.

4. How does the iterative soft thresholding algorithm (ISTA) work for LASSO?

ISTA iteratively updates the weight vector using the following update rule:

$$w_i = S_{\lambda\alpha\gamma} \left(w_{i-1} - \frac{2\gamma}{n} X_n^T (Y_n - X_n w_{i-1}) \right)$$

Where $S_{\lambda\alpha\gamma}$ is the soft-threshold operator that shrinks some weights to zero.

5. Compare Matching Pursuit (MP) and Orthogonal Matching Pursuit (OMP).

- MP: Greedily selects the features that correlates the most with the residual and updates weights sequentially.

- OMP: Instead of updating the weight sequentially, it orthogonalizes the selected features to improve accuracy.

6. When should Elastic Net regularization be used instead of LASSO?

Elastic Net should be used when the features are highly correlated, as LASSO tends to select one feature arbitrarily. Elastic Net balances l_1 -norm sparsity with l_2 -norm stability, preventing over selection of a single features.

7. Given the optimization problem for Elastic Net, what does the hyperparameter α control?

$$\min_{w \in \mathbb{R}^D} \frac{1}{n} \sum_{i=1}^n (y_i - f_w(x_i))^2 + \lambda(\alpha |w|_1 + (1 - \alpha) |w|_2^2)$$

- $\alpha = 1$: LASSO behavior (sparse solution)
 - $\alpha = 0$: Ridge Regression (No sparsity, only regularization)
 - $0 < \alpha < 1$ A balance between sparsity and stability
8. What is the main difference between Tikhonov regularization (Ridge) and LASSO?
- Ridge regression (l_2 -norm) shrinks all weights but doesn't enforce sparsity (no feature selection).
 - LASSO (l_1 -norm) forces some weights to be exactly zero, leading to feature selection

9. What are the key steps in the Matching Pursuit (MP) algorithm?

- Initialize residual $r_0 = Y_n$, coefficient vector $w = 0$.
- Select the variable the correlates the most with the residual.
- Update the coefficient vector for selected variable.
- Update the residual.
- Repeat until a stopping criterion is met.

10. What is the main limitation of LASSO when handling highly correlated features?

LASSO randomly selects one feature from a group of correlated features while ignoring the rest, potentially losing important information. Elastic Net solves this by incorporating both l_1 and l_2 penalties.